


Dahlia

Predictable Accelerator Design with
Time-Sensitive Affine types

Rachit Nigam, Sachille Atapattu, Ted Bauer, Adrian Sampson, Zhiru Zhang

What is an FPGA?

FIELD -
PROGRAMMABLE
GATE
ARRAY

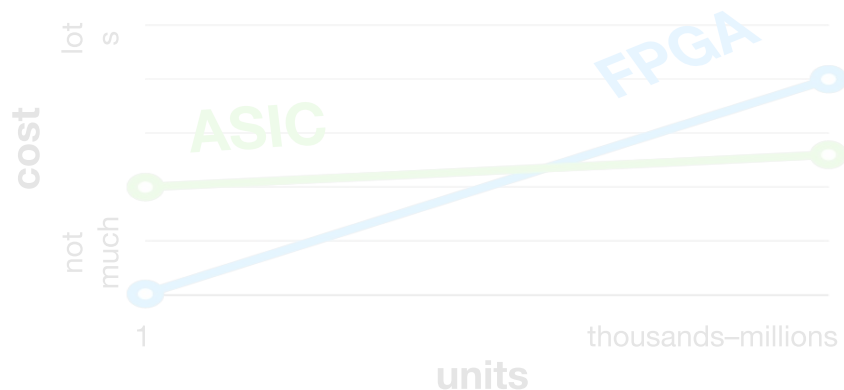


What is an FPGA?

FIELD-PROGRAMMABLE
GATE
ARRAY

A circuit emulator.

To prototype or replace an ASIC.




signal processing
networking
embedded devices

VLSI validation
radio
microcontrollers


A general accelerator.

To offload computation from a CPU.



machine learning
signal processing
search

graph analytics
genomics
scientific computing



Amazon

EC2 F1 instances let you rent FPGAs by the hour, in the cloud.



Microsoft

Project Catapult: accelerate Bing search.
Project Brainwave: low-latency DNN inference.

A Reconfigurable Fabric for Accelerating Large-Scale Datacenter Services

Andrew Putnam¹ Adrian M. Caulfield¹ Eric S. Chung¹ Derek Chiou¹
Kypros Constantinides² John Demme³ Hadi Esmaeilzadeh⁴ Jeremy Fowers⁵
Gopi Prashanth Gopal¹ Jan Gray¹ Michael Haselman¹ Scott Hauck⁵ Stephen Heil¹
Amir Hormati⁶ Joo-Young Kim¹ Sitaram Lanka¹ James Larus⁷ Eric Peterson¹
Simon Pope¹ Aaron Smith¹ Jason Thong¹ Phillip Yi Xiao¹ Doug Burger¹

Microsoft

Abstract

Datacenter workloads demand high computational capabilities, flexibility, power efficiency, and low cost. It is challenging to improve all of these factors simultaneously. To advance datacenter capabilities beyond what commodity server designs can provide, we have designed and built a composable, reconfigurable fabric to accelerate portions of large-scale software services. Each instantiation of the fabric consists of a 6x8 2-D torus of high-end Stratix V FPGAs embedded into a half-rack of 48 machines. One FPGA is placed into each server, accessible through PCIe, and wired directly to other FPGAs with pairs of 10 Gb SAS cables.


In this paper, we describe a medium-scale deployment of this fabric on a bed of 1,632 servers, and measure its efficacy in accelerating the Bing web search engine. We describe the requirements and architecture of the system, detail the

desirable to reduce management issues and to provide a consistent platform that applications can rely on. Second, datacenter workloads evolve extremely rapidly, making non-programmable hardware features impractical. Thus, datacenter providers are faced with a conundrum: they need continued improvements in performance and efficiency, but cannot obtain those improvements from general-purpose systems.

Reconfigurable chips, such as Field Programmable Gate Arrays (FPGAs), offer the potential for flexible acceleration of many workloads. However, as of this writing, FPGAs have not been widely deployed as compute accelerators in either datacenter infrastructure or in client devices. One challenge traditionally associated with FPGAs is the need to fit the accelerated function into the available reconfigurable area. One could virtualize the FPGA by reconfiguring it at run-time to support more functions than could fit into a single device. However, current reconfiguration times for standard FPGAs


23
AUTHORS!

A circuit emulator.



≠

A general accelerator.



RTL

register-transfer level

Verilog





VHDL

Bluespec

Chisel



Actual Silicon



C



```
graph TD; C[C] --> RTL[RTL<br/>register-transfer level];
```

The diagram illustrates the High-Level Synthesis (HLS) process. It starts with a gray box labeled 'C' at the top left. A gray arrow points downwards from this box to another gray box below it. This second box is labeled 'RTL' in large bold letters, with 'register-transfer level' written in smaller text underneath. To the right of the arrow, the text 'High-Level Synthesis' is written in a large, bold, black font.

High-Level Synthesis

RTL

register-transfer level

image and video processing, financial analytics, bioinformatics, and scientific computing applications. Since RTL programming in VHDL or Verilog is unacceptable to most application software developers, it is essential to provide a highly automated compilation/synthesis flow from C/C++ to FPGAs.

As a result a growing number of FPGA designs are

Pitfalls of C-based HLS

HLS tools *transparently* compile C/C++ code and
predictably generate performant hardware *automatically*.

```
void bias(float* vec, float* biases);
```

```
void bias(float* vec, float* biases);
```

Unsupported memory access on the variable 'biases', which is (or contains) an array with unknown size at compile time.

```
int fac(int n, int acc) {  
    if (n == 0)  
        return acc;  
    else  
        return fac(n - 1, n * acc);  
}
```

```
int fac(int n) {  
    if (n == 0)  
        return 1;  
    else  
        return n * fac(n - 1);  
}
```

```
int fac(int n, int acc) {  
    if (n == 0)  
        return acc;  
    else  
        return fac(n - 1, n * acc);  
}
```



```
int fac(int n) {  
    if (n == 0)  
        return 1;  
    else  
        return n * fac(n - 1);  
}
```




HLS tools *transparently* compile C/C++ code and *predictably generate* performant hardware *automatically*.


- An *undefined subset* of C with simple control and static operators supported.

HLS tools *transparently* compile C/C++ code and *predictably generate* performant hardware *automatically*.


- An *undefined subset* of C with simple control and static operators supported.

```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```





```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    #pragma HLS UNROLL factor=5  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```



```
#pragma HLS ARRAY_PARTITION variable=A factor=5  
#pragma HLS ARRAY_PARTITION variable=B factor=5
```

```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    #pragma HLS UNROLL factor=5  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```



```
#pragma HLS ARRAY_PARTITION variable=A factor=5  
#pragma HLS ARRAY_PARTITION variable=B factor=5
```

```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    #pragma HLS UNROLL factor=5  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```

```
graph LR; i[i] --> Ai[A[i]]; i --> Bi[B[i]]; factor5[factor=5] --> Ai; factor5 --> Bi;
```

HLS tools *transparently* compile C/C++ code and *predictably generate* performant hardware *automatically*.


- An *undefined subset* of C with simple control and static operators supported.
- *Second class* primitives for controlling design decisions. User hints needed for high performance design.


HLS tools ~~transparently~~ compile C/C++ code and ~~predictably generate~~ performant hardware ~~automatically~~.

- An *undefined subset* of C with simple control and static operators supported.
- *Second class* primitives for controlling design decisions. User hints needed for high performance design.

```
#pragma HLS ARRAY_PARTITION variable=A factor=5  
#pragma HLS ARRAY_PARTITION variable=B factor=5
```

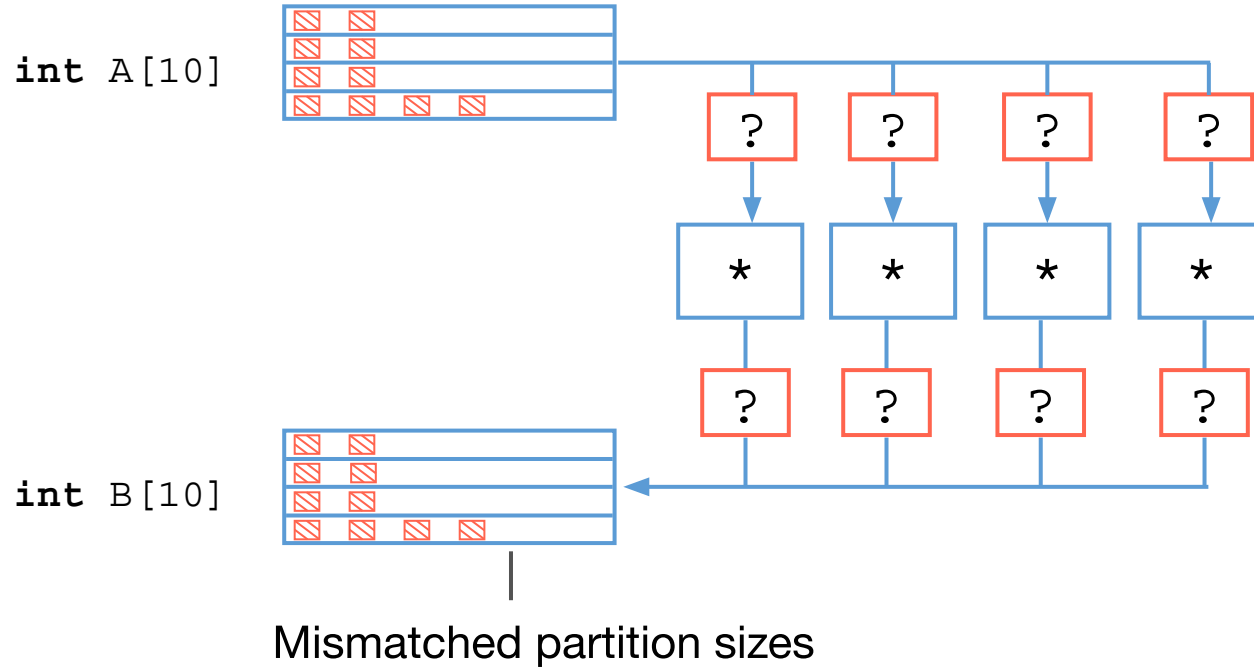
```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    #pragma HLS UNROLL factor=4  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```






```
#pragma HLS ARRAY_PARTITION variable=A factor=4  
#pragma HLS ARRAY_PARTITION variable=B factor=4
```

```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    #pragma HLS UNROLL factor=4  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```





```
#pragma HLS ARRAY_PARTITION variable=A factor=5
#pragma HLS ARRAY_PARTITION variable=B factor=5
```

```
int A[10];
int B[10];
for (int i = 0; i < 10; i++) {
    #pragma HLS UNROLL factor=5
    int x = A[i];
    int y = x * 5;
    B[i] = y;
}
```



```
#pragma HLS ARRAY_PARTITION variable=A factor=4
#pragma HLS ARRAY_PARTITION variable=B factor=4
```

```
int A[10];
int B[10];
for (int i = 0; i < 10; i++) {
    #pragma HLS UNROLL factor=4
    int x = A[i];
    int y = x * 5;
    B[i] = y;
}
```



```
#pragma HLS ARRAY_PARTITION variable=A factor=M  
#pragma HLS ARRAY_PARTITION variable=B factor=M
```

```
int A[10];  
int B[10];  
for (int i = 0; i < 10; i++) {  
    #pragma HLS UNROLL factor=M  
    int x = A[i];  
    int y = x * 5;  
    B[i] = y;  
}
```

HLS tools ~~transparently~~ compile C/C++ code and ~~predictably generate~~ performant hardware ~~automatically~~.

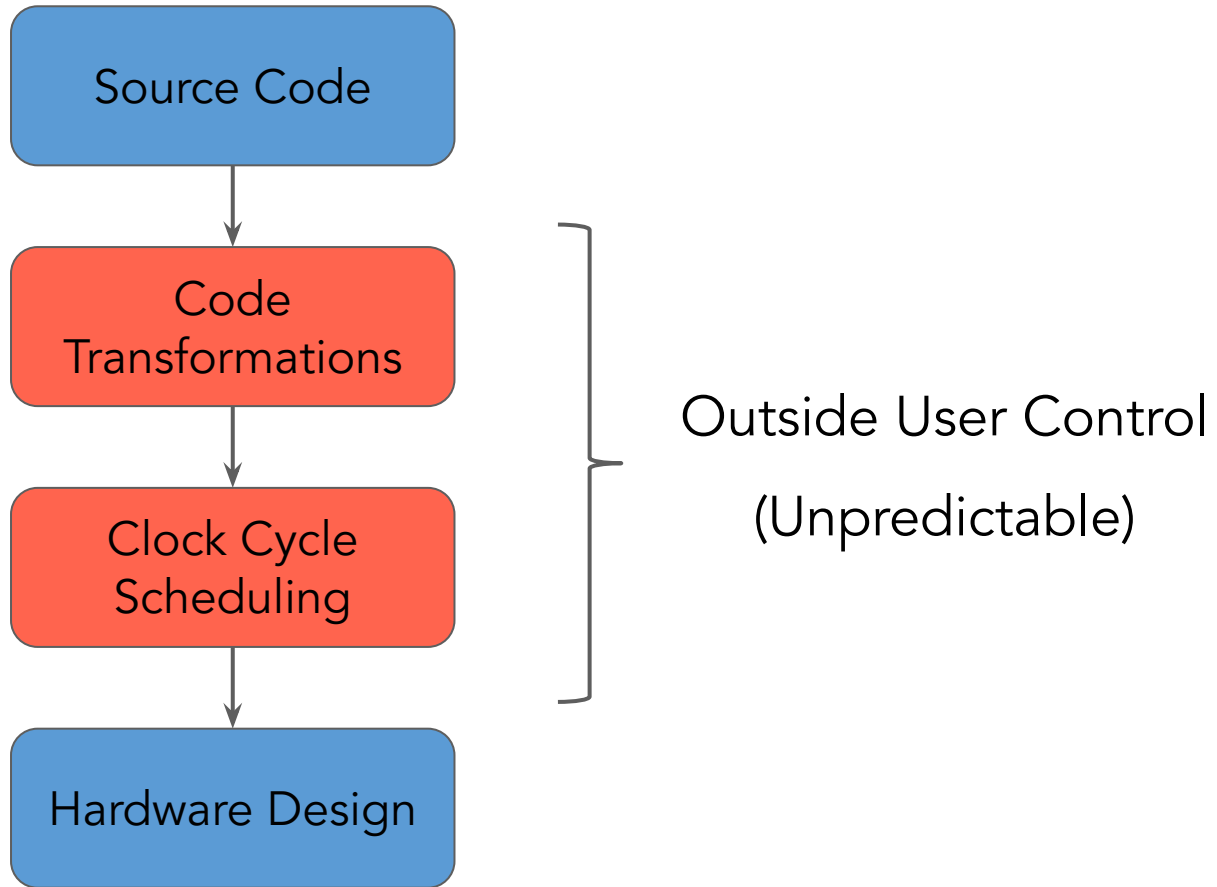
- An *undefined subset* of C with simple control and static operators supported.
- *Second class* primitives for controlling design decisions. User hints needed for high performance design.
- Non trivial additions to designs to support parallel semantics for C.

HLS tools ~~transparently~~ compile C/C++ code and ~~predictably generate~~ performant hardware ~~automatically~~.

- An *undefined subset* of C with simple control and static operators supported.
- *Second class* primitives for controlling design decisions. User hints needed for high performance design.
- Non trivial additions to designs to support parallel semantics for C.

HLS tools ~~transparently~~ compile C/C++ code and
~~predictably generate~~ performant hardware ~~automatically~~.

HLS tools compile an *ill-defined*, imperative language and *unpredictably generate* performant designs with *user aid*.

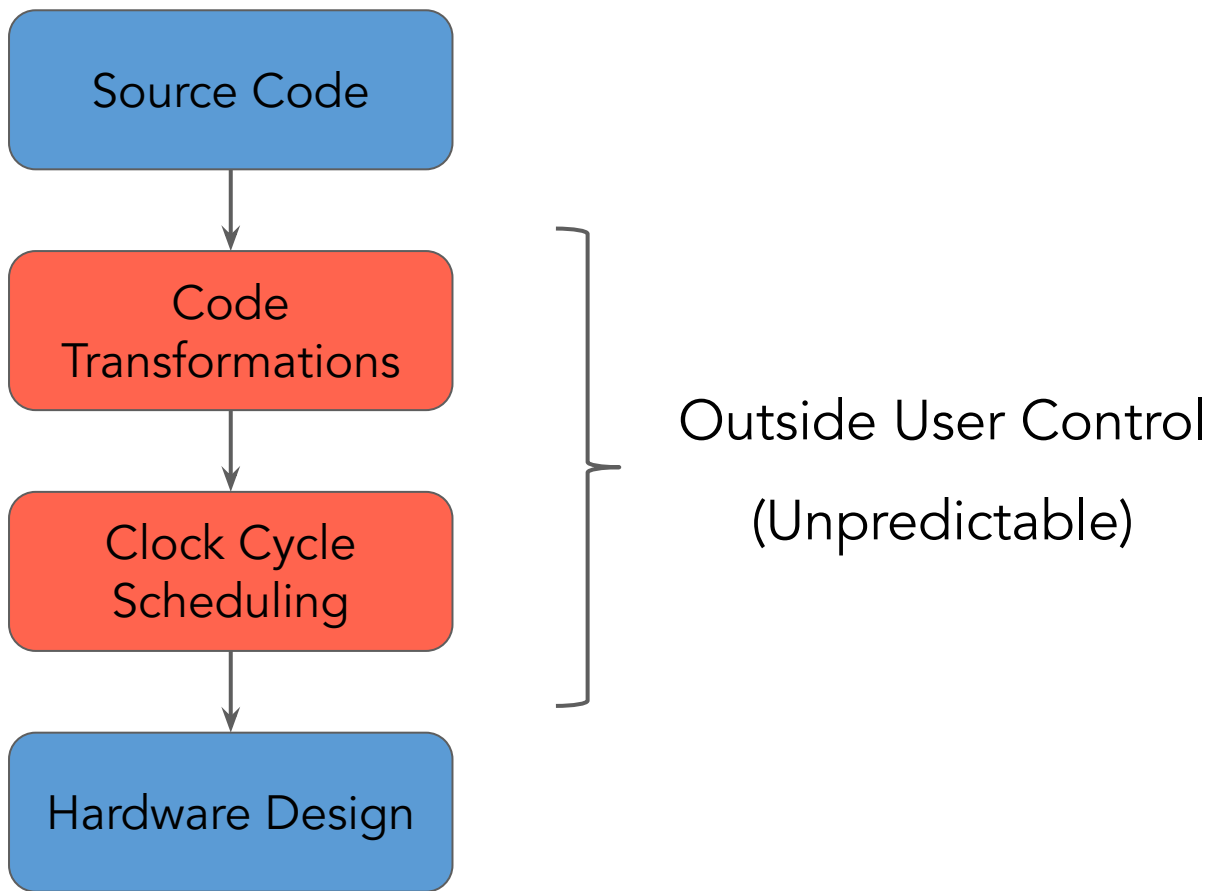


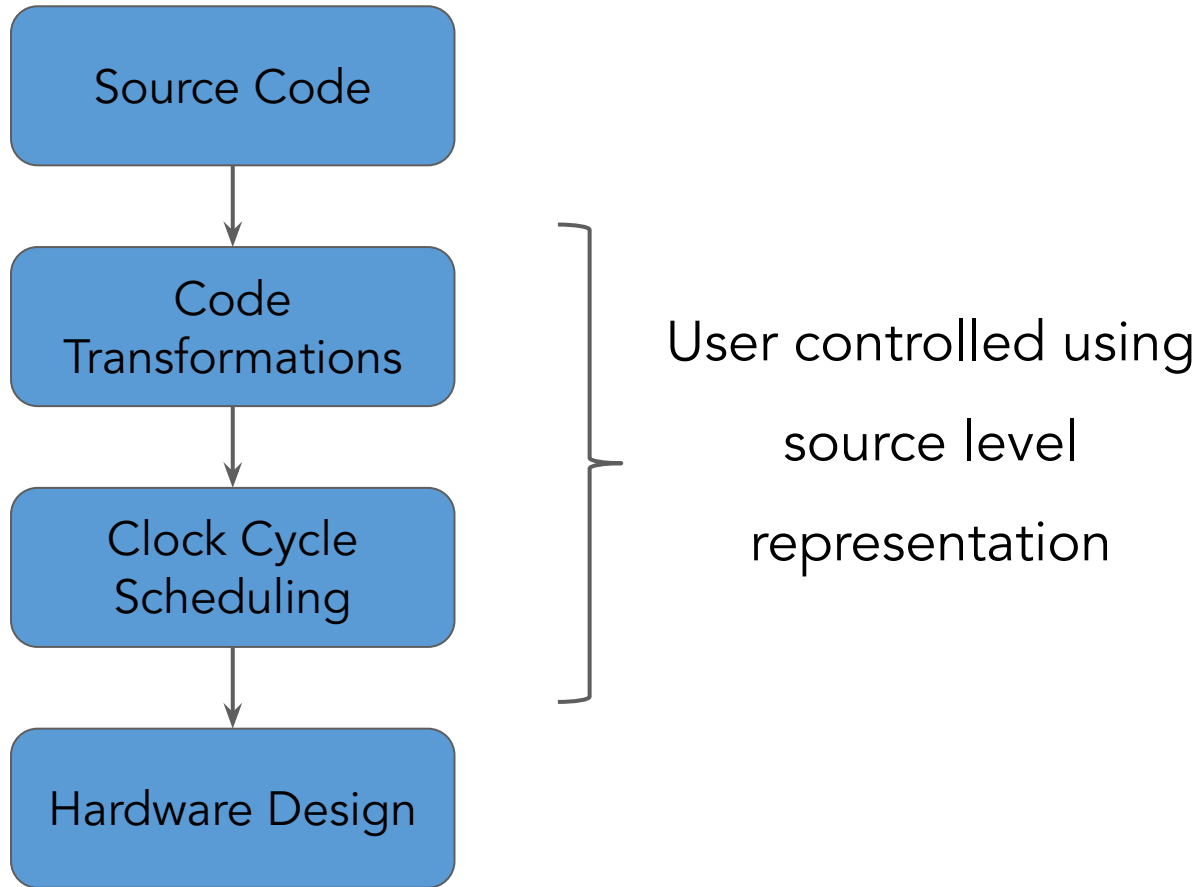
Dahlia

Compiles a *well-defined* imperative language
and *predictably* generates performant hardware
with *user-aid*

Predictable Hardware Generation

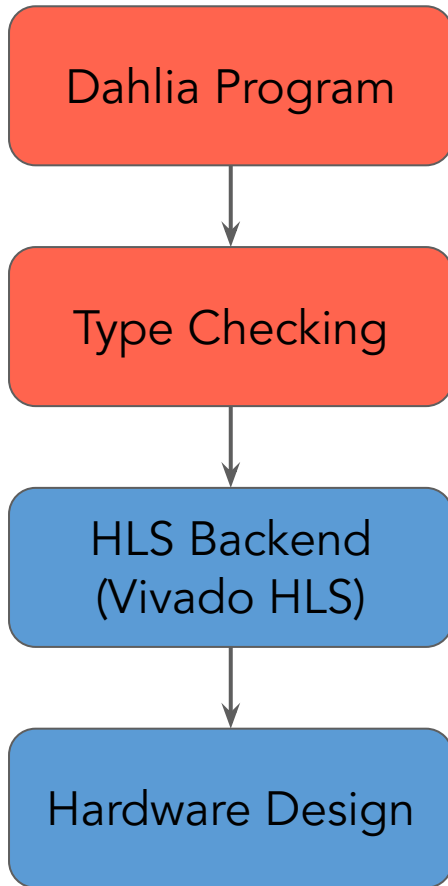
- Only allow programs that represent *hardware cost* at the source level.
- Reject programs that require complicated transformations to work in hardware.






Design Goals

- *Source-level reasoning* over implicit scheduling decisions.
- *Predictable hardware generation*: Use an affine type system to disallow programs with unpredictable hardware generation.
- *First class scheduling primitives* to control design decisions. Guarantee that scheduling primitives don't change semantics.



```
decl A: float[10];
```

```
decl A: float[10];
```



Global memories.
Mapped to BRAMs.

```
decl A: float[10];
```

```
let x = A[0];
```

```
A[9] = 1;
```



```
decl A: float[10];
```

```
let x = A[0];
```

```
A[9] = 1;
```



Error: Physical resource `A' already used in this context

A ; B

A and B *may* run in *parallel*

```
let x = A[0];
```

```
A[9] = 1;
```

Conflict! Cannot read and write to `A` in the same cycle. Assume all memories are single ported.

```
decl A: float[10];
```

```
let x = A[0]
```

```
A[9] = 1;
```



A logical time step

A - - - B

Run A and then run B.
Create a *logical timestep*.

```
let x = A[0]
```

```
---
```

```
A[9] = 1
```

No conflict.
Statements run in separate timesteps.

$\{ A \quad --- \quad B \} ;$
 $\{ C \quad --- \quad D \} ;$
E

Specify complex parallel behavior using ---
and ;

```

decl A: bit<32>[M][N];
decl x: bit<32>[N];
decl y: bit<32>[N];
decl tmp: bit<32>[M];

for (let i = 0..N) {
  y[i] := 0.0;
}
---
for (let i = 0..M) {
  tmp[i] := 0.0;
  ---
  for (let j = 0..N) {
    let t = tmp[i];
    ---
    tmp[i] := t + A[i][j] * x[j];
  }
  ---
  for (let j = 0..N) {
    let y0 = y[j];
    ---
    y[j] := y0 + A[i][j] * tmp[i];
  }
}

```

Source level
reasoning for cycle
boundaries.


```


decl A: bit<32>[M] [N];
decl x: bit<32>[N];
decl y: bit<32>[N];
decl tmp: bit<32>[M];

for (let i = 0..N) {
  y[i] := 0.0;
}
---
for (let i = 0..M) {
  tmp[i] := 0.0;
  ---
  for (let j = 0..N) {
    let t = tmp[i];
    ---
    tmp[i] := t + A[i][j] * x[j];
  }
  ---
  for (let j = 0..N) {
    let y0 = y[j];
    ---
    y[j] := y0 + A[i][j] * tmp[i];
  }
}

```

Source level
reasoning for cycle
boundaries.

```
decl A: bit<32>[10];  
decl B: bit<32>[10];  
for (i = 0 to 10) {  
    A[i] = A[i + 1] - 1;  
}
```




Bounded *DOALL* loops. Cross iteration dependencies disallowed.

```
decl A: bit<32>[10];  
decl B: bit<32>[10];  
for (i = 0 to 10) {  
    A[i] = A[i + 1] - 1;  
}
```



Error! Cross iteration
dependency.

```
decl A: bit<32>[10];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
}
```



First class unrolling primitive.
Typechecker reasons about array
access exclusion with unrolling.

```
decl A: bit<32>[10];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
}
```




Error: Insufficient resources for parallel access.
Required: 5, Available: 1.

```
decl A: bit<32>[10 bank 5];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
}
```

First class banking primitive.
Generates extra resources for
parallel accesses.

```
decl A: bit<32>[10 bank 5];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
} combine {  
    B[0] *= x  
}
```




Explicit instantiation for reduction trees. Can contain arbitrary computation pipelines.

```

decl A: bit<32>[10 bank 5];
decl B: bit<32>[10];
for (i = 0 to 10) unroll 5 {
    let x = A[i] + 1;
} combine {
    B[0] *= x
}


```




```

decl A: bit<32>[10 bank 5];
decl B: bit<32>[10];
for (i = 0 to 10) unroll 5 {
    let x = A[i] + 1;
} combine {
    B[0] *= x
}

```




```
decl A: bit<32>[10 bank 10];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
} combine {  
    B[0] *= x  
}
```

```
decl A: bit<32>[10 bank 10];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
} combine {  
    B[0] *= x  
}
```

Error! Banking factor and unrolling factor do not match for access `A[i]`


```
decl A: bit<32>[10 bank 5];  
decl B: bit<32>[10];  
for (i = 0 to 10) unroll 5 {  
    let x = A[i] + 1;  
} combine {  
    B[0] *= x  
}
```



```

decl A: bit<32>[10 bank 10];
decl B: bit<32>[10];
for (i = 0 to 10) unroll 5 {
    let x = A[i] + 1;
} combine {
    B[0] *= x
}

```



```

decl A: bit<32>[10 bank 10];
decl B: bit<32>[10];

view sh_A = shrink A[by 5];

for (...) unroll 5 {
    let x = sh_A[i] + 1;
} combine {
    B[0] *= x
}

```

Memory View : Generate additional hardware to access `A` as a 5 banked memory.

```
decl A: bit<32>[10 bank 10];  
decl B: bit<32>[10];  
  
view sh_A = shrink A[by 5];  
  
for (...) unroll 5 {  
    let x = sh_A[i] + 1;  
} combine {  
    B[0] *= x  
}
```

Type checking ensures that parallelization/pipelining directives preserve *functional correctness* of the sequential program.

```

decl A: bit<32>[10 bank 10];
decl B: bit<32>[10];

view sh_A = shrink A[by 5];

for (...) unroll 5 {
    let x = sh_A[i] + 1;
} combine {
    B[0] *= x
}

```

Type checking ensures that parallelization/pipelining directives preserve *functional correctness* of the sequential program.



Dahlia

- *Source-level reasoning* over implicit scheduling decisions.
- *Predictable hardware generation*: Use an affine type system to disallow programs with unpredictable hardware generation.
- *First class scheduling primitives* to control design decisions. Guarantee that scheduling primitives don't change semantics.

Evaluation

- Machsuite benchmarks. Port 18 HLS C programs to Dahlia.
- Initial minimal effort porting: Syntactically rewrite C programs to Dahlia.
- Design Space Exploration with Dahlia: Rewrite ports to use Dahlia's parallelization primitives. Compare with DSE in HLS C.

Evaluation



Future Work

- *Affine functions*: Area-efficiency trade-offs with reusable hardware modules.
- *Multi-ported memories*: Crucial for high performance designs on modern FPGAs.
- *Constraint guided DSE*: Use type constraints to generate constraints for DSE.

Thanks!

capra.cs.cornell.edu/fuse

The screenshot shows the GitHub repository page for `cucapra / dahlia`. The repository has 4 issues, 4 pull requests, 0 projects, 0 wiki pages, 1 security issue, 1 insight, and 1 setting. The description states: "Affine types for predictable hardware generation" with a link to <https://capra.cs.cornell.edu/fuse>. The repository has 1,412 commits, 9 branches, 1 release, 1 environment, 8 contributors, and is licensed under MIT. The latest commit is by `rachitnigam` with the message "rebrand" 21 hours ago. The file list includes:

File	Description	Last Commit
<code>.circleci</code>	try mozilla/sbt on circleci	last month
<code>docs</code>	predictable hardware	6 days ago
<code>notes</code>	warning about notes	2 months ago
<code>project</code>	add project assembly dep	7 months ago
<code>src</code>	2.13 migration fix	21 hours ago
<code>tools</code>	else is a keyword for vim	21 hours ago
<code>website</code>	expand on overview. Set it as default landing page	2 months ago
<code>.gitignore</code>	gitignore	2 months ago
<code>.hook.yaml</code>	cleanup main makefile and remove seashell.js bundle	11 months ago
<code>LICENSE</code>	Add a license	11 months ago
<code>Makefile</code>	Create new docs website.	6 months ago
<code>README.md</code>	rebrand	21 hours ago
<code>build.sbt</code>	upgrade to scala 2.13	21 hours ago
<code>fuse</code>	link to scala 2.13 jar for fuse executable	21 hours ago

Constrain directed Optimization

Constrain directed Optimization

```
#pragma HLS ARRAY_PARTITION variable=A factor=M
#pragma HLS ARRAY_PARTITION variable=B factor=L

int A[10];
int B[10];
for (int i = 0; i < K; i++) {
    #pragma HLS UNROLL factor=N
    int x = A[i];
    int y = x * 5;
    B[i] = y;
}
```

Constrain directed Optimization

```
#pragma HLS ARRAY_PARTITION variable=A factor=M
#pragma HLS ARRAY_PARTITION variable=B factor=L


int A[10];
int B[10];
for (int i = 0; i < K; i++) {
    #pragma HLS UNROLL factor=N
    int x = A[i];
    int y = x * 5;
    B[i] = y;
}
```

$M == L$

$M \% N == 0$

$K \% N == 0$

Constrain directed Optimization




Fuse



RTL


Register-Transfer Level



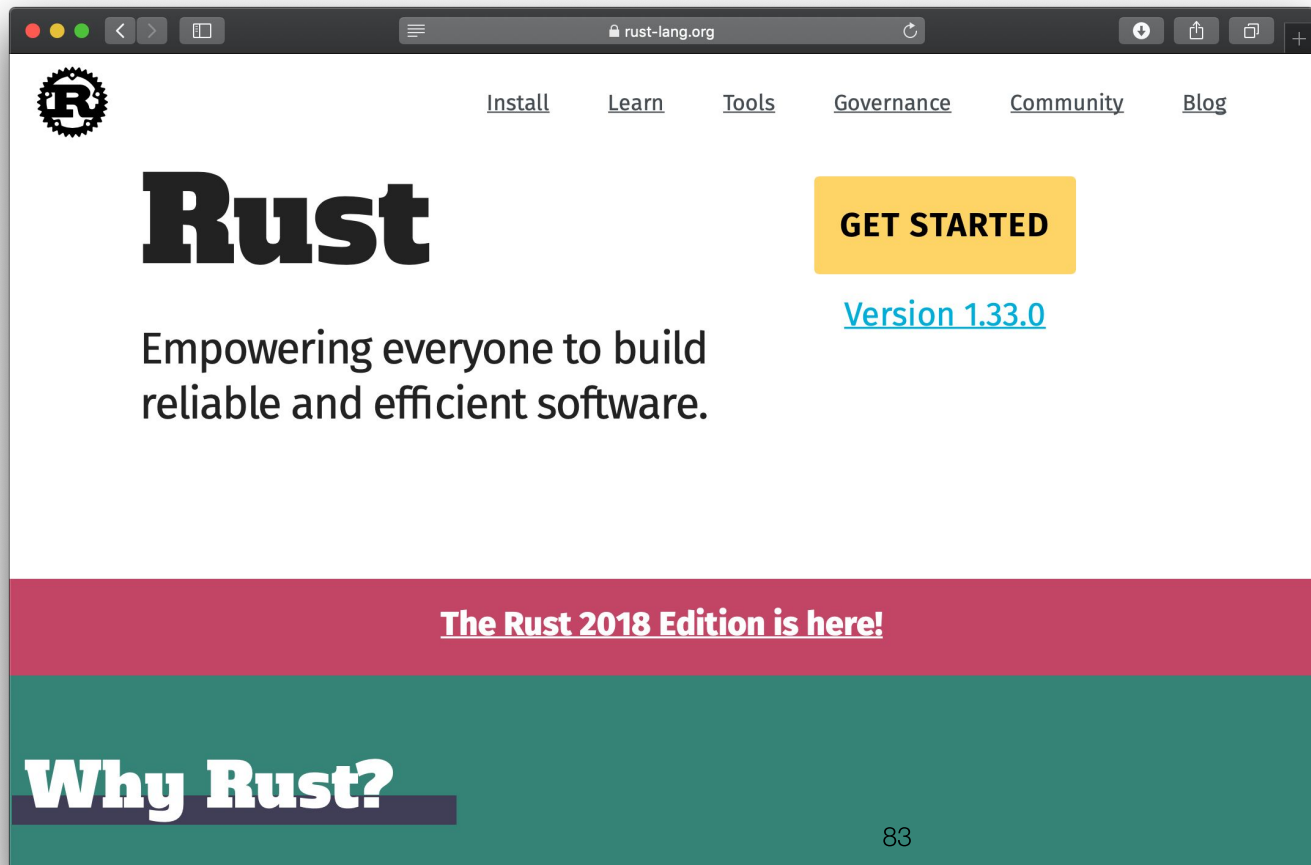
C



Assembly



Affine types and **linear types**, as made famous recently by **Rust**.





Banked memory types

```
memory bank(5) A : int[10];
```

A[0] A[5]
A[1] A[6]
A[2] A[7]
A[3] A[8]
A[4] A[9]

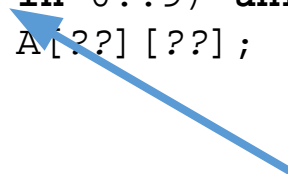
Banked memory types

```
memory bank(5) A : int[10];  
for (let i in 0..1) {  
  access A[0][i];  
  access A[1][i];  
  access A[2][i];  
  access A[3][i];  
  access A[4][i];  
}  
  
// one access to each  
// A[j] allowed here  
  
STATIC  DYNAMIC 
```

A[0][0] A[0][1]
A[1][0] A[1][1]
A[2][0] A[2][1]
A[3][0] A[3][1]
A[4][0] A[4][1]

Hybrid indices for unrolling

```
memory bank(5) A : int[10];  
for (let i in 0..9) unroll 5 {  
    access A[??][??];  
}
```



Hybrid indices for unrolling

```
memory bank(5) A : int[10];  
for (let i in 0..9) unroll 5 {  
    access A[i];  
}
```